Ray Tracing

ab-gh

1 December 2021

Task: To create an advanced optical simulation program in C++. Using the principles of *Ray Tracing*, *Object-Oriented Design*, and the material covered in PH3170, this project aims to produce a realistic image of a simulated scene, including accurate objects and lighting.

1   Introduction and Theory

To start the project, I set out a series of objectives for the final code, and a set of milestones to progress through. In order to fulfill the requirements of the task, my code needs to be able to:

- Allow a user to specify the geometry, location, and material properties of primative shapes in the scene.

- Allow a user to specify the location and intensity of sources of light.

- Simulate the object's interactions with the light sources and other objects in the scene using *reverse ray tracing*.

- Produce an image of the simulated scene, as if a camera was taking a photograph.

To achieve these aims, I set out to complete the following milestones, roughly in order. By breaking the project down into smaller components, I not only ensured that I didn't waste time working on one feature when an underlying feature was broken, but also applied my code-writing process to Object Oriented Design (OOD) principles.

1. Construct a fully-functional but abstract Vector data type, capable of defining vectors and points in 3D cartesian space, and computing all relevant vector operations.

2. Construct a "virtual camera", which sends rays out into the "world".

3. Define primative objects, classes which contain the object's geometry and properties, as well as functions to calculate the intersections between itself and a camera ray, and a function to calculate the surface normal of the shape at any point.

4. Define light sources, with an intensity and position.

5. Construct a "rendering engine", capable of simulating various types of illumination, and linking the virtual camera to the "world" (scene) of objects in front of it.

1.1 Ray Tracing Theory

To produce a realistic depiction of the scene, this project uses a technique called *reverse ray tracing.* In a real physical scene, the illumination of objects can be approximated as light sources emitting rays which interact with the object, reflecting light to a new ray direction. If one of these light rays, after any number of reflections, hits a sensor (be it an eye or camera), it is "seen" by the observer. In a computational simulation, a tiny fraction of the rays emitted from a source will actually make it to the observer, so this process is very inefficient.

To combat this, *reverse* ray tracing is used, with the same principles, except the observer emits a "sight" ray, which reflects according to the same laws, until it hits a light source. To produce the image, the observer "looks" through a viewport - a rectangular boundary in space - similar to a camera sensor. Sight rays are constructed with a vector from an "eye" (focus of the observation plane) to each pixel on the viewport, which is then projected into the scene, and the reverse ray tracing begins.
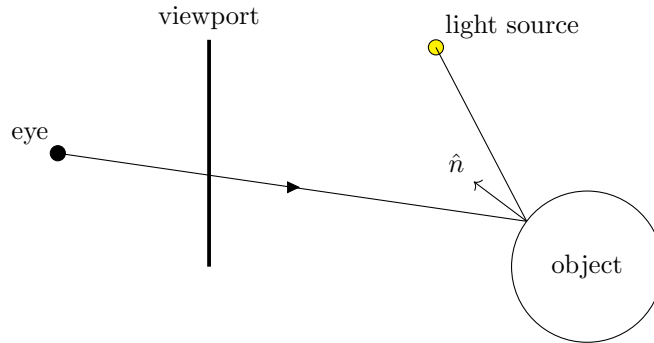

Figure 1: Demonstration of reverse ray tracing, reflecting a sight ray from the eye through the viewport off a sphere.

1.2 Object Oriented Design

This project required a lot of different components to work together accurately, yet still abstract enough to be flexible, allowing the development of more shapes, interactions, and materials in future. Because of this, applying the principles of Object Oriented Design was essential. Each component of the project was "compartmentalised" in its own header file, using inheritance where possible. Each class interacts with the renderer, which produces the final image.

2 Implementation

This section explains the implementation of the ray tracing simulator, following the process outline in the milestones above.

2.1 Vector base class

The first task of the implementation was to construct a fully-featured vector class. Defined in `Vector.h`, the `Vector` class serves as the basis for almost every other element of the simulator. As

such, it defines an extensive list of methods and operations for vectors. Operations for vectors are also defined outside of the `Vector` class, as inline functions, so that they can be accessed outside of a Vector. In addition to the basic mathematical operations, the header file defines the dot and cross product (`dot(u,v)`, `cross(u,v)`), unit vector (`unit(u)`), as well as reflections (`reflect(i,n)`) and refractions (`refract(i,n,eta)`). The `Vector` class also has three aliases, which help to make their specific use clearer in context: `V3`, `P3`, `RGB`.

RGB vectors have their own class definition, `RGB.h`. These force the values to be "clamped" between 0 and 255 when output, and provide a `write_RGB` function to write an RGB value out to a file.

2.2 Camera, viewport, and scene

Now that the basic vector class is implemented, the virtual world of the simulation can be built. The file `Camera.h` defines the `Camera` class, which includes the eye and viewport. For simplicity, this has default values of a 16:9 aspect ratio image, with 1000 pixels across, which translates to a height of 2 in the simulated world. The focal length is defined as 1, which means the "eye" is one unit behind the centre of the viewport. The viewport is centered on the origin, looking down the negative z axis. This geometry allows you to consider the image as if it were an x, y plane, with depth as $-z$.

Next, the `Scene.h` file defines the simulated world as the class `Scene`. This is a very simple class, with two `std::vector` attributes: one to hold the list of objects in the scene, and one to hold the list of light sources. Two methods - `addObject()` and `addSource()` are also defined, which add objects and light sources to the scene once created.

`Scene.h` also defines `save()` and `read()` methods, which save the existing scene to a "scene" file, or load in a generated scene file. These allow for complex scenes to be defined without creating a complex `main.cpp` file, or allow a "base" world (such as a basic room with soft lighting) to be imported, and extra objects added (as in the demonstration `main.cpp`).

`Ray.h` defines the aforementioned "sight"-rays. It is a very simple class, defining a ray as having a vector origin and unit vector direction. It also defines a `()` operator, which allows the calculation of a ray at "time" t or "distance" μ much like one would define a function: `Ray(x)`.

2.3 Objects, intersections, and normals

With a class now defined to hold objects in a world, the next step was to define those objects. To render an object in the world, each object class requires a series of defining parameters (which can define a unique instance of the object), and two methods: one to calculate where a vector line will intersect with the object's surface, and one to calculate the normal to the object's surface at any point. If these two functions are provided, the surface of any object can be detected and rendered by the renderer.

`Object.h` defines several different objects to choose from, as well as an abstract base class used to build the other objects. `Object` class defines the functions needed for an object: a parameter constructor, a string constructor (used to deserialize information when calling `Scene.read()`), an intersection function, a normal function, and a serializer (for `Scene.save()`). One feature of note is the return type of the `Object.intersect()` function. This function returns a pair `std::pair<double, const Object*>`, where the double is the unit direction multiplier of the ray's intersection point. While it may seem redundant to return the object that has been intersected, this allows complex objects to be created with "sub-objects". For instance, the `Cone` object is comprised of an "empty" cone, with a base cap, which is actually a `Disc` sub-object.

This project currently implements the following shapes:

- `Sphere`
    - `Vector` position
    - `double` radius
    - `Vector` colour
    - `double` reflectivity

- `InfinitePlane`
    - `Vector` position
    - `Vector` normal
    - `Vector` colour
    - `double` reflectivity

- `Disc`
    - `Vector` position
    - `Vector` normal
    - `double` radius
    - `Vector` colour
    - `double` reflectivity

- `Cone`
    - `Vector` position
    - `Vector` base
    - `double` radius
    - `Vector` colour
    - `double` reflectivity

2.4   Sources and intensity

Now that objects can be defined, light sources need to be provided to allow them to be illuminated. These are defined in `Source.h`, which defines a `Source` class. This requires a position and intensity, and provides a deserialization constructor and a serializer. While only invisible point-sources are defined, it would be easy to extend this class to create visible sources (i.e. emissive objects) with intersection and normal functions, or different light sources such as a "sun" (i.e. a source at infinity with parallel light rays) or a directional light (i.e. a source which only provides illumination at some solid angle).

2.5   Rendering and Image Generation

The final - and by far the most complex - requirement for the project is the "rendering engine". This is defined in `Render.h` and, while not a class, defines a multitude of functions. Figure 2 shows a simple flowchart describing the operation of the rendering engine.

When the `render()` function is called, and is passed a valid scene and camera, it will loop over each pixel in the viewport and construct a ray from the "eye" to that pixel. It then projects the ray into the scene, and follows a process to determine the pixel's colour:

1. For each object in the scene's object list, find where the ray intersects the object.

2. For the closest object, calculate the intensity by summing, for each source:

   - Diffuse illumination, by taking the dot product of the vector from the intersection point to the source, with the surface normal at the intersection point.

   - Specular illumination, by taking the dot product of the vector from the intersection point to the source, with the reflected ray around the surface normal at the intersection point.

   - Reflective illumination, by taking the reflected ray around the surface normal at the intersection point, and treating that as a new sight ray, projecting it into the scene, and going to step 2.
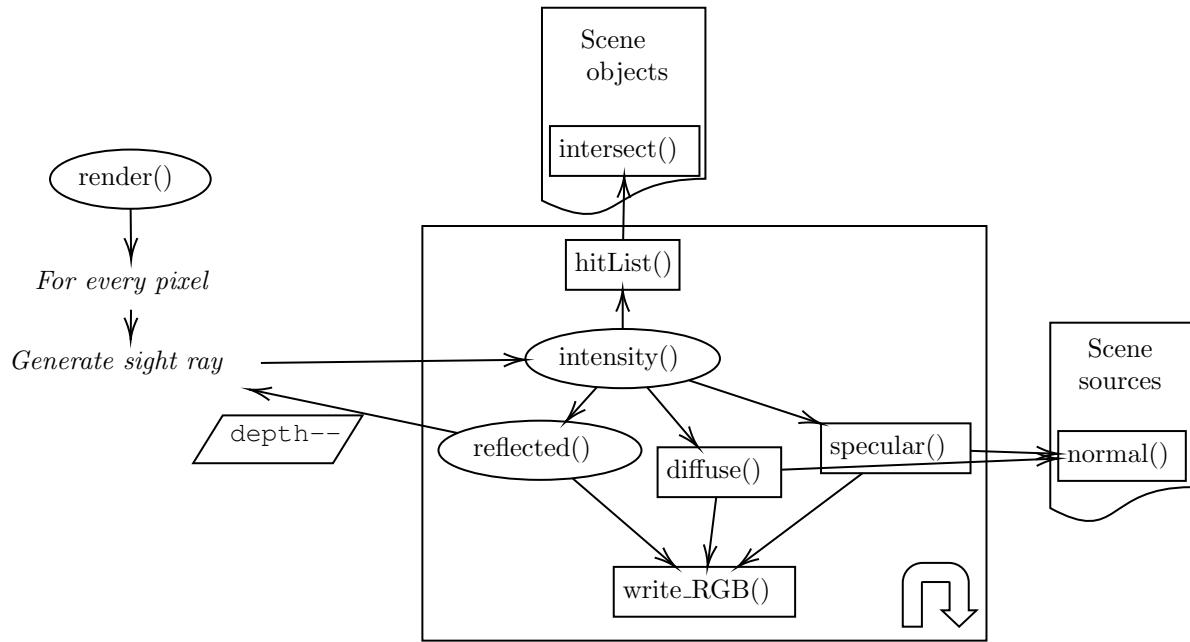
Figure 2: Flowchart describing the operation of the rendering engine.

3. Write the pixel's RGB values out to the image file.

To generate each "sight" ray, the renderer creates vectors from the lower-left corner of the viewport along the x and y axis to the pixel. As the renderer also knows the vector from the origin to the eye, and the origin to the lower-left corner, it can construct a vector from the eye to the pixel. That vector is then used to define the sight ray's direction, and the eye is used as the origin. This is shown below in Figure 3.
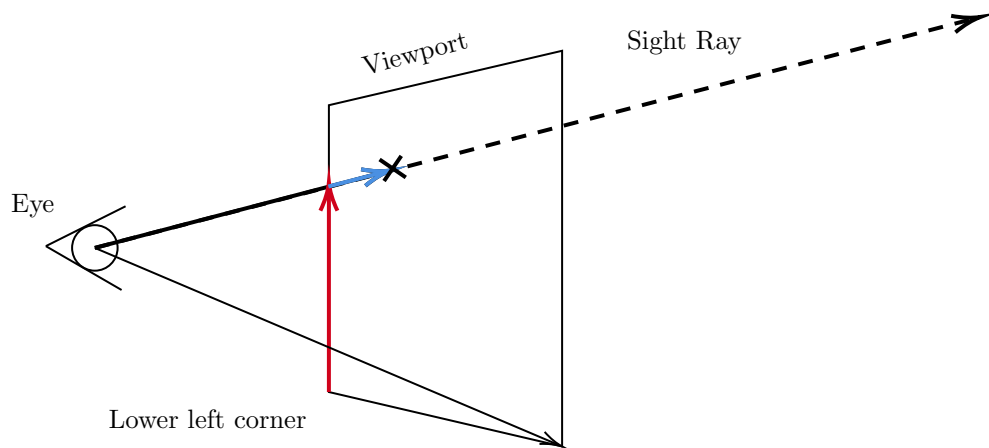

Figure 3: Diagram showing how each sight ray is constructed.

5

2.6   Main and demonstration

The final piece of the puzzle is the contents of `main.cpp`. To successfully construct a scene and output an image, the user needs to include all the headers, create a camera and a scene, add objects to the scene or load a scene file, and call the `render(camera, scene, depth)` function. The default `main.cpp` file included with this project will generate a scene demonstrating all of the features of the project, by loading in a basic empty room, adding two spheres, a cone, and a source. The code may take a few minutes to run, and will output the image as `~cmake-build-debug/image.ppm`, but hopefully the output will be worthwhile:
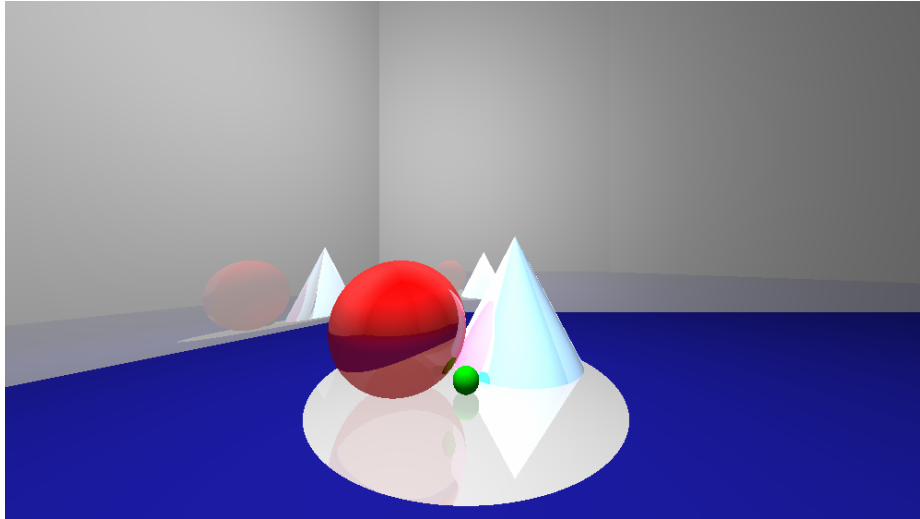

Figure 4: Demonstration image of a ray-tracing scene, produced by this project.

3   Conclusion and Future Ideas

In conclusion, I believe that this project fulfills or succeeds the requirements set out by the script. By closely following OOD methodology, the project is easily broken down and understood, and is flexible and can be easily updated. As I enjoyed this project greatly, I have a few ideas of features and improvements I'd like to implement in the future:

- Allow parameters for the initialization of the `Camera` class, allowing for a "free camera" to be positioned anywhere in the scene.

- Creating visible light sources, and emissive objects, and other types of light sources (such as directional, sunlight or softbox).

- Creating transparency/translucency parameters for objects, and implementing ray refraction. The `Vector.refract` method already exists for this, but is currently unused.

- Adding shadows, depth of field, and other features mentioned in the `README.md` roadmap.
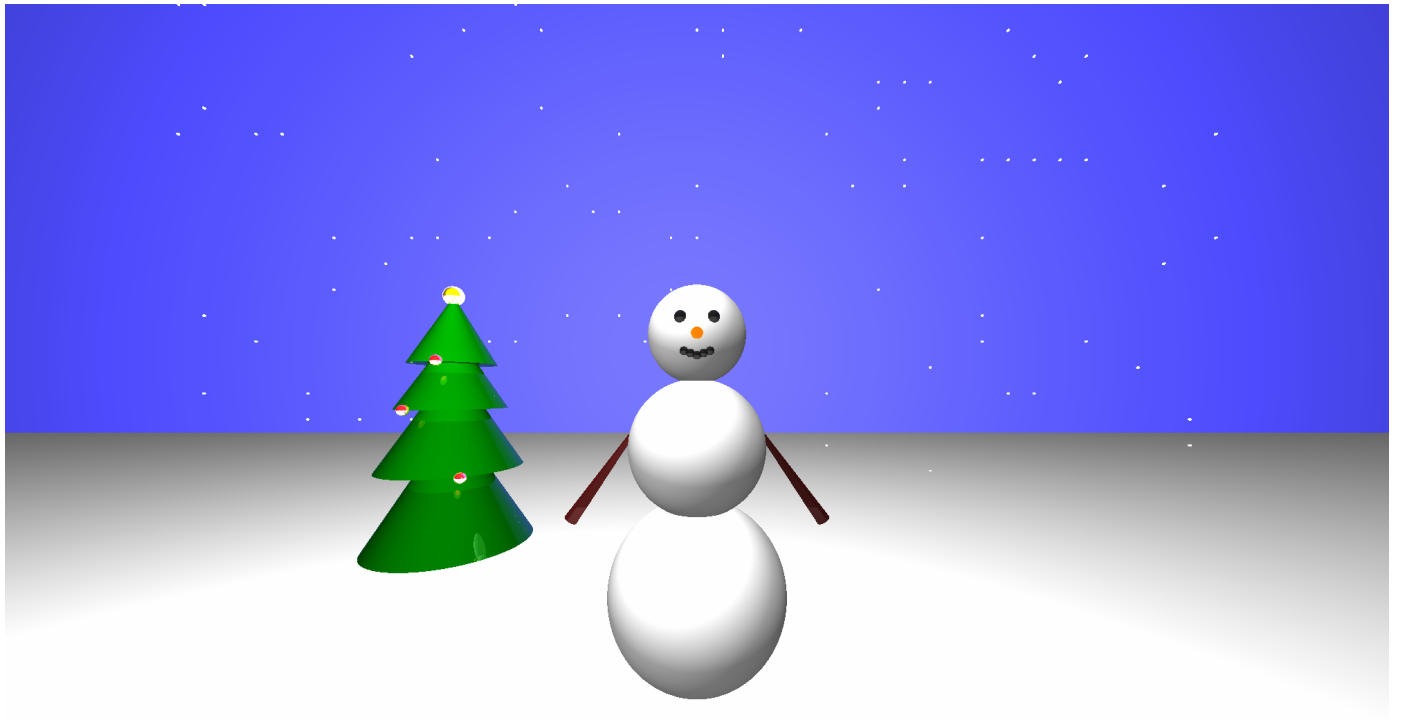
A   One more thing...

Figure 5: Merry Christmas and a Happy New Year! (Image code provided in *snowman.cpp*).